

# Ad Impression Counter - Test Assignment

Create a concurrent service that tracks ad impressions across multiple campaigns in real-time, with mechanisms to avoid counting duplicate impressions from the same user within a specified time period.

*Note: Consider the task as a guideline. You have a great freedom to experiment, to showcase interesting snippets, ideas, workarounds.*

## Task

The goal is to build an in-memory impression counter that supports concurrent processing, including handling duplicate impressions efficiently and providing a REST API to manage campaigns, track impressions, and retrieve statistics.

## Instructions

### 1. Implement Core Data Structures:

- **Campaign:** Represents an advertising campaign.
- **Impression:** Represents each ad view, including **userID** and **timestamp**.
- **Stats:** Stores aggregated impression statistics for each campaign.

### 2. Duplicate Handling:

- Track impressions uniquely for each **userID** to avoid duplicate counting.
- Implement a TTL (Time-To-Live) mechanism to discard duplicate records after a set time period (e.g., one hour).
- Ensure that each unique impression from a user is only counted once within the specified TTL.

### 3. Concurrency:

- Use concurrency controls (e.g., channels, mutexes) to handle multiple impressions tracked simultaneously without race conditions.
- Make the impression tracking system thread-safe.

### 4. Build REST API Endpoints:

- **POST /api/v1/campaigns** — Register a new campaign.
  - Request Body: **CreateCampaignRequest**
- **POST /api/v1/impressions** — Track a new impression.
  - Request Body: **TrackImpressionRequest**
- **GET /api/v1/campaigns/{id}/stats** — Get impression statistics for a specific campaign.
  - Response Body: **Stats**

## Data Structures

// Campaign represents an advertising campaign

```
type Campaign struct {
    ID          string    `json:"id"`
    Name        string    `json:"name"`
    StartTime   time.Time `json:"start_time"`
}
```

// Impression represents a single ad view

```
type Impression struct {
    CampaignID string    `json:"campaign_id"`
    Timestamp   time.Time `json:"timestamp"`
    UserID      string    `json:"user_id"`
    AdID        string    `json:"ad_id"`
}
```

// Stats represents aggregated impression statistics

```
type Stats struct {
    CampaignID string `json:"campaign_id"`
    LastHour   int64 `json:"last_hour"`
    LastDay    int64 `json:"last_day"`
    TotalCount int64 `json:"total"`
}
```

## API Requests/Responses

```
type CreateCampaignRequest struct {
    Name        string    `json:"name"`
    StartTime   time.Time `json:"start_time"`
}
```

```
type TrackImpressionRequest struct {
    CampaignID string `json:"campaign_id"`
    UserID      string `json:"user_id"`
    AdID        string `json:"ad_id"`
}
```

// Error response format

```
type ErrorResponse struct {  
    Error string `json:"error"`  
    Code  int    `json:"code"`  
}
```

## Technical Focus

- **Concurrency:** Use channels and mutexes to handle concurrent updates safely.
- **Duplicate Check:** Implement a TTL for `userID` records to prevent duplicate counting within a specified time period (e.g., one hour).
- **API Design:** Structure endpoints for intuitive use and handle edge cases.
- **Project Structure:** Organize code with scalability in mind, allowing for future growth.
- **Minimum Dependencies:** Prefer standard, or well-known packages.

## Example Workflow

1. **Register a Campaign:**
  - Send a `POST` request to `/api/v1/campaigns` with `CreateCampaignRequest` in the JSON body.
  - Returns the new campaign's ID and a success message.
2. **Track Impressions:**
  - Send a `POST` request to `/api/v1/impressions` with `TrackImpressionRequest` in the body.
  - Check for duplicate impressions (based on `userID` and TTL).
  - Returns a success message.
3. **Get Campaign Statistics:**
  - Send a `GET` request to `/api/v1/campaigns/{id}/stats`.
  - Returns the aggregated `Stats` for the campaign, including `LastHour`, `LastDay`, and `TotalCount`.

## Evaluation Criteria

1. **Concurrency Handling:**
  - Are channels and mutexes used effectively to manage concurrent updates and prevent race conditions?
  - Is the TTL for duplicate check implemented correctly?
2. **API Design:**
  - Are the endpoints structured in a RESTful manner?
  - Is duplicate detection efficient and scalable?
3. **Code Organization and Readability:**
  - Is the project organized and structured to support future expansion?
  - Is the code clear and well-commented?

**4. Error Handling and Documentation:**

- Are error cases handled appropriately, with clear messages?
- Is the code and API well-documented?

**5. Testing:**

- Include tests to validate core functionality (e.g., concurrent impressions and duplicate check).